

Project Report on California Housing Price Prediction



**Submitted by
Debatri Roy
PGP DS MAR 2021 cohort 1**

Python execution O/P Screenshots and Interpretation

■ I have done this project in two parts:

- Data analysis and cleaning as explained in **EDA and data cleaning**.
- Training of machine learning models explained in **Training Machine Learning Algorithms**.

EDA and Data Cleaning

- Creating new features
- Removing outliers
- Transforming skewed features
- Checking for multicollinearity

Training machine learning algorithms

- Linear Regression
- Ridge Regression
- Support Vector Regression
- Gradient Boosting Regression
- Stacking of various models

(Contd.)

Python execution O/P Screenshots and Interpretation

■ Including library functions of packages

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: import numpy.random as rnd
from statsmodels.stats.outliers_influence import variance_inflation_factor
import statsmodels.api as sm
```

```
In [3]: from sklearn.base import BaseEstimator,TransformerMixin,RegressorMixin
```

```
In [4]: rnd.seed(42)
```


(Contd.)

Python execution O/P Screenshots and Interpretation

■ Load Dataset

```
In [5]: # Load dataset
housing = pd.read_excel("1553768847_housing.xlsx")
housing
```

Out[5]:



	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	ocean_proximity	median_house_value
0	-122.23	37.88	41	880	129.0	322	126	8.3252	NEAR BAY	452600
1	-122.22	37.86	21	7099	1106.0	2401	1138	8.3014	NEAR BAY	358500
2	-122.24	37.85	52	1467	190.0	496	177	7.2574	NEAR BAY	352100
3	-122.25	37.85	52	1274	235.0	558	219	5.6431	NEAR BAY	341300
4	-122.25	37.85	52	1627	280.0	565	259	3.8462	NEAR BAY	342200
...
20635	-121.09	39.48	25	1665	374.0	845	330	1.5603	INLAND	78100
20636	-121.21	39.49	18	697	150.0	356	114	2.5568	INLAND	77100
20637	-121.22	39.43	17	2254	485.0	1007	433	1.7000	INLAND	92300
20638	-121.32	39.43	18	1860	409.0	741	349	1.8672	INLAND	84700
20639	-121.24	39.37	16	2785	616.0	1387	530	2.3886	INLAND	89400

20640 rows x 10 columns

We can see that the dataset contains 20640 rows and 10 columns.


(Contd.)

Python execution O/P Screenshots and Interpretation

■ Print first few rows of this data

```
In [6]: housing.head()
```

Out[6]:



	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	ocean_proximity	median_house_value
0	-122.23	37.88	41	880	129.0	322	126	8.3252	NEAR BAY	452600
1	-122.22	37.86	21	7099	1106.0	2401	1138	8.3014	NEAR BAY	358500
2	-122.24	37.85	52	1467	190.0	496	177	7.2574	NEAR BAY	352100
3	-122.25	37.85	52	1274	235.0	558	219	5.6431	NEAR BAY	341300
4	-122.25	37.85	52	1627	280.0	565	259	3.8462	NEAR BAY	342200

The head() displays the first 5 rows of housing dataset.


(Contd.)

Python execution O/P Screenshots and Interpretation

■ First Portion: EDA and Data Cleaning

● Getting information of Data Frame

```
In [7]: housing.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  int64  
3   total_rooms            20640 non-null  int64  
4   total_bedrooms        20433 non-null  float64
5   population             20640 non-null  int64  
6   households             20640 non-null  int64  
7   median_income          20640 non-null  float64
8   ocean_proximity        20640 non-null  object  
9   median_house_value     20640 non-null  int64  
dtypes: float64(4), int64(5), object(1)
memory usage: 1.6+ MB
```

The info() gives the information of the “housing” data frame like no. of non-null values, data type and memory usage.

(Contd.)

Python execution O/P Screenshots and Interpretation

● Checking for missing values

```
In [8]: # checking for missing values
housing.isnull().sum()
```



```
Out[8]: longitude           0
latitude                   0
housing_median_age         0
total_rooms                0
total_bedrooms            207
population                 0
households                 0
median_income              0
ocean_proximity            0
median_house_value         0
dtype: int64
```

Here, we can see in the output that, only “total_bedrooms” contains null values and the total number of null values are 207.

(Contd.)

Python execution O/P Screenshots and Interpretation

● Ocean Proximity is a categorical variable. Let's see what values it contains.

```
In [9]: housing['ocean_proximity'].value_counts()
```



```
Out[9]: <1H OCEAN      9136  
        INLAND        6551  
        NEAR OCEAN    2658  
        NEAR BAY      2290  
        ISLAND         5  
        Name: ocean_proximity, dtype: int64
```

The above output shows, ISLAND has only 5 entries. I'll be adding these to other class in subsequent stages (in GetDummies class).

(Contd.)

Python execution O/P Screenshots and Interpretation

● Getting description of Data Frame

```
In [10]: housing.describe()
```



Out[10]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value
count	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000	20640.000000	20640.000000	20640.000000	20640.000000
mean	-119.569704	35.631861	28.639486	2635.763081	537.870553	1425.476744	499.539680	3.870671	206855.816909
std	2.003532	2.135952	12.585558	2181.615252	421.385070	1132.462122	382.329753	1.899822	115395.615874
min	-124.350000	32.540000	1.000000	2.000000	1.000000	3.000000	1.000000	0.499900	14999.000000
25%	-121.800000	33.930000	18.000000	1447.750000	296.000000	787.000000	280.000000	2.563400	119600.000000
50%	-118.490000	34.260000	29.000000	2127.000000	435.000000	1166.000000	409.000000	3.534800	179700.000000
75%	-118.010000	37.710000	37.000000	3148.000000	647.000000	1725.000000	605.000000	4.743250	264725.000000
max	-114.310000	41.950000	52.000000	39320.000000	6445.000000	35682.000000	6082.000000	15.000100	500001.000000

The `describe()` shows some basic statistical details like count, mean, median, standard deviation, quartile values, percentile, minimum and maximum values of each attributes of the dataset.

(Contd.)




Python execution O/P Screenshots and Interpretation

Histogram of housing Data Frame

```
In [11]: housing.hist(bins=50,figsize=(15,10))
```

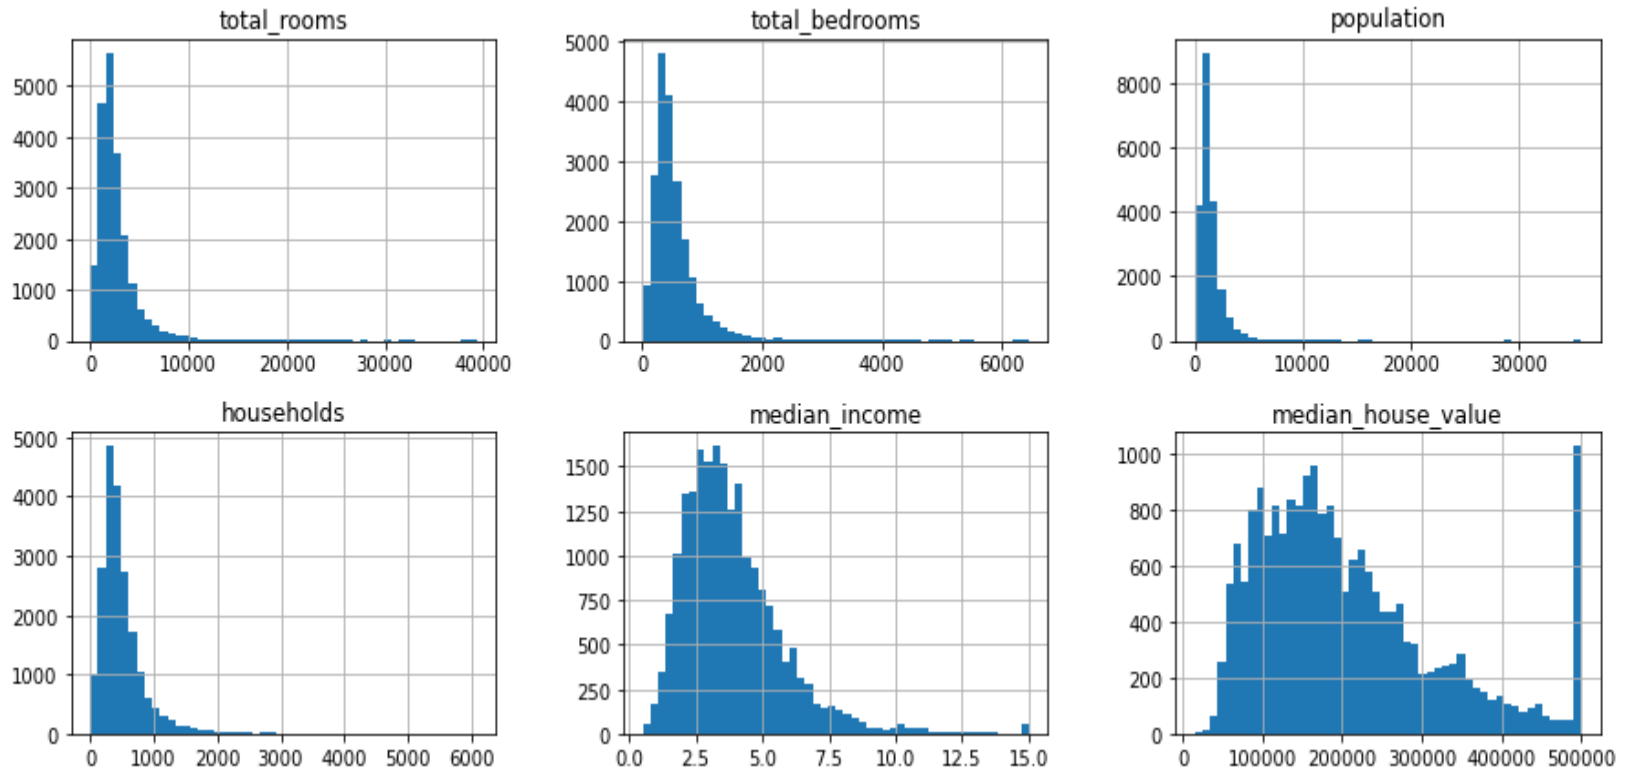


```
Out[11]: array([[<AxesSubplot:title={'center':'longitude'}>,  
                <AxesSubplot:title={'center':'latitude'}>,  
                <AxesSubplot:title={'center':'housing_median_age'}>],  
               [<AxesSubplot:title={'center':'total_rooms'}>,  
                <AxesSubplot:title={'center':'total_bedrooms'}>,  
                <AxesSubplot:title={'center':'population'}>],  
               [<AxesSubplot:title={'center':'households'}>,  
                <AxesSubplot:title={'center':'median_income'}>,  
                <AxesSubplot:title={'center':'median_house_value'}>]],  
            dtype=object)
```



(Contd.)

Python execution O/P Screenshots and Interpretation



From above histograms it's clear that most of the features are skewed. It seems that 'housing_median_age' and 'median_house_value' have been capped at the end (Peak at the end).

(Contd.)

Python execution O/P Screenshots and Interpretation

● Getting Test Dataset:



Used Stratified sampling technique.



Defined a new feature **income_cat** which is **income category** and used it for sampling.

```
In [12]: housing['income_cat']=np.ceil(housing['median_income']/1.5)
housing['income_cat'].where(housing['income_cat']<5,5.0,inplace=True)
```

```
In [13]: from sklearn.model_selection import StratifiedShuffleSplit
```

```
In [14]: split=StratifiedShuffleSplit(n_splits=1,test_size=0.2,random_state=42)
for train_index,test_index in split.split(housing,housing['income_cat']):
    strat_train_set=housing.loc[train_index]
    strat_test_set=housing.loc[test_index]
```

```
In [15]: strat_train_set.drop('income_cat',axis=1,inplace=True)
strat_test_set.drop('income_cat',axis=1,inplace=True)
```

```
In [16]: strat_train_set.to_csv("strat_train_set.csv",index=False)
strat_test_set.to_csv("strat_test_set.csv",index=False)
```


(Contd.)

Python execution O/P Screenshots and Interpretation

Exploratory Data Analysis

```
In [17]: data=pd.read_csv('strat_train_set.csv')
```

```
In [18]: data.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 16512 entries, 0 to 16511
Data columns (total 10 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   longitude            16512 non-null  float64
1   latitude             16512 non-null  float64
2   housing_median_age   16512 non-null  int64
3   total_rooms          16512 non-null  int64
4   total_bedrooms       16354 non-null  float64
5   population           16512 non-null  int64
6   households           16512 non-null  int64
7   median_income        16512 non-null  float64
8   ocean_proximity      16512 non-null  object
9   median_house_value   16512 non-null  int64
dtypes: float64(4), int64(5), object(1)
memory usage: 1.3+ MB
```


The info() gives the information of the “data” data frame like no. of non-null values, data type and memory usage.

(Contd.)

Python execution O/P Screenshots and Interpretation

```
In [19]: plt.figure(figsize=(10,6))  
plt.scatter(x=data['longitude'],y=data['latitude'])  
plt.title("Distribution of households",size=16)
```

Out[19]: Text(0.5, 1.0, 'Distribution of households')



The above plot showing the distribution of households according to longitude and latitude values in California.

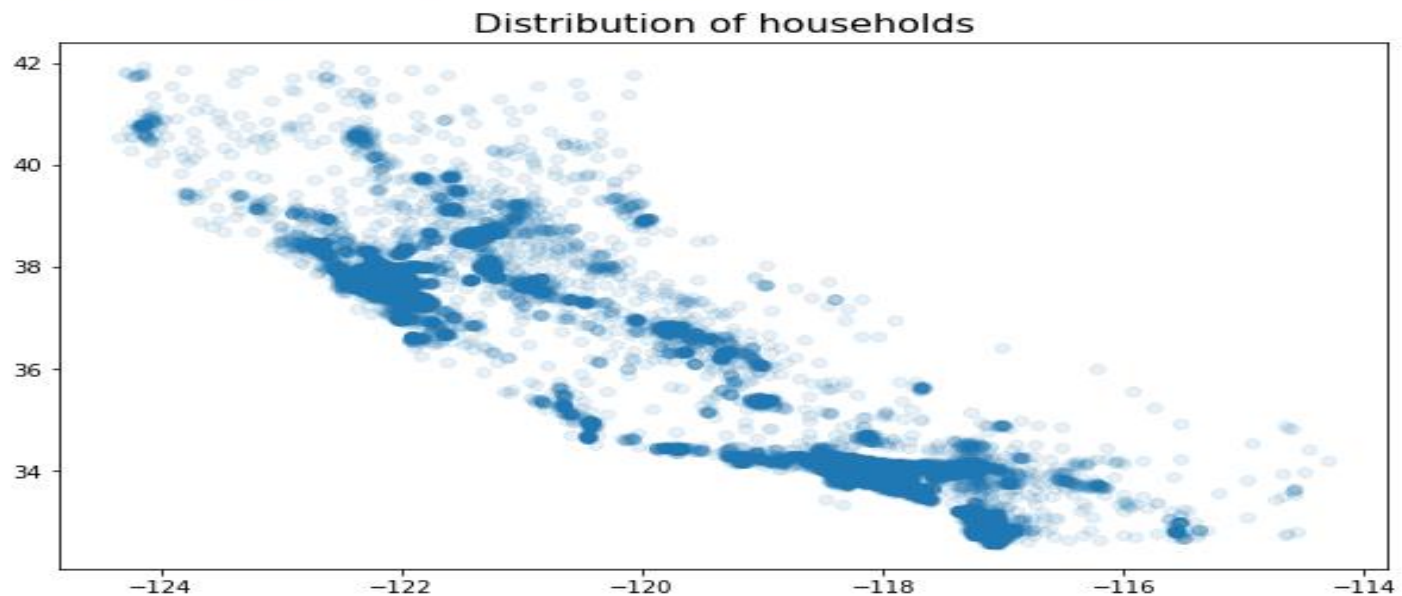
(Contd.)

Python execution O/P Screenshots and Interpretation

```
In [20]: plt.figure(figsize=(10,6))  
plt.scatter(x=data['longitude'],y=data['latitude'],alpha=0.1)  
plt.title("Distribution of households",size=16)
```



```
Out[20]: Text(0.5, 1.0, 'Distribution of households')
```



By setting alpha = 0.1, we can see high density areas in the above plot.

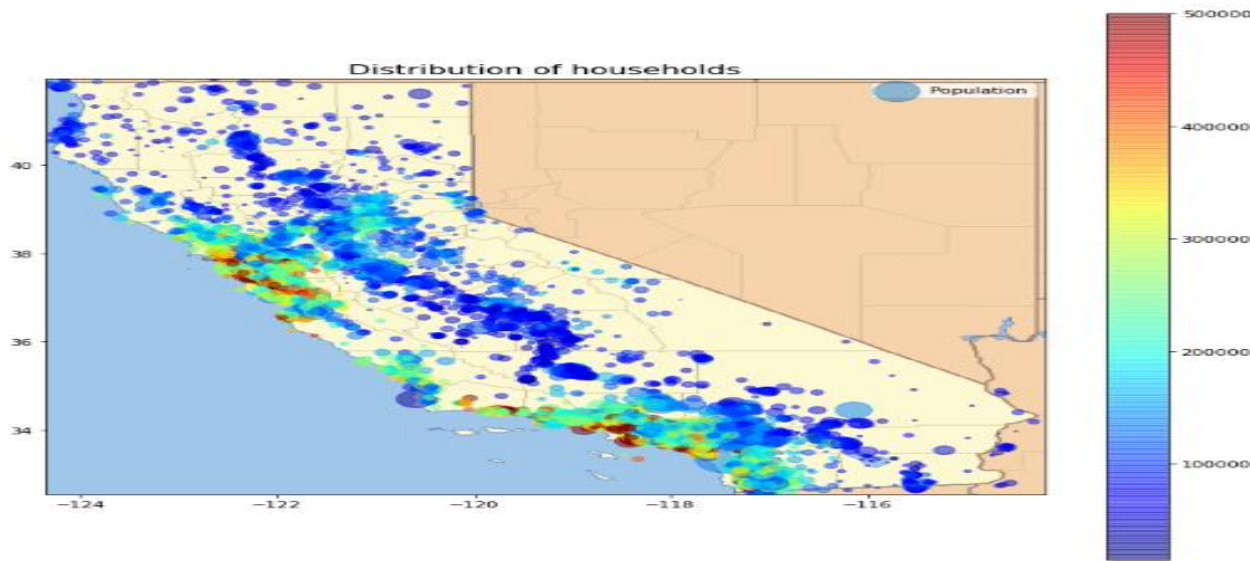
(Contd.)

Python execution O/P Screenshots and Interpretation

```
In [21]: plt.figure(figsize=(12,12))
img=plt.imread('california.png')
plt.imshow(img,zorder=0,extent=[-124.35,-114.2,32.54,41.95])

plt.scatter(x=data['longitude'],y=data['latitude'],alpha=0.5,s=data['population']/30,c=data['median_house_value'],
            cmap=plt.get_cmap("jet"),zorder=1,label='Population')
plt.colorbar()
plt.title("Distribution of households",size=16)
plt.legend()
```

Out[21]: <matplotlib.legend.Legend at 0xc123d90>



From the above plot we can infer that,

- Housing prices are much related to location and population density.
- Housing prices near ocean are higher except in northern California.

(Contd.)

Python execution O/P Screenshots and Interpretation



Now, let's see the correlation of '**median house value**' with other columns. This is **Pearson's correlation coefficient**.

```
In [22]: corr_matrix=data.corr()
```

```
In [23]: corr_matrix['median_house_value'].sort_values(ascending=False)
```



```
Out[23]: median_house_value    1.000000  
         median_income        0.687160  
         total_rooms          0.135097  
         housing_median_age    0.114110  
         households           0.064506  
         total_bedrooms        0.047689  
         population           -0.026920  
         longitude            -0.047432  
         latitude             -0.142724  
         Name: median_house_value, dtype: float64
```

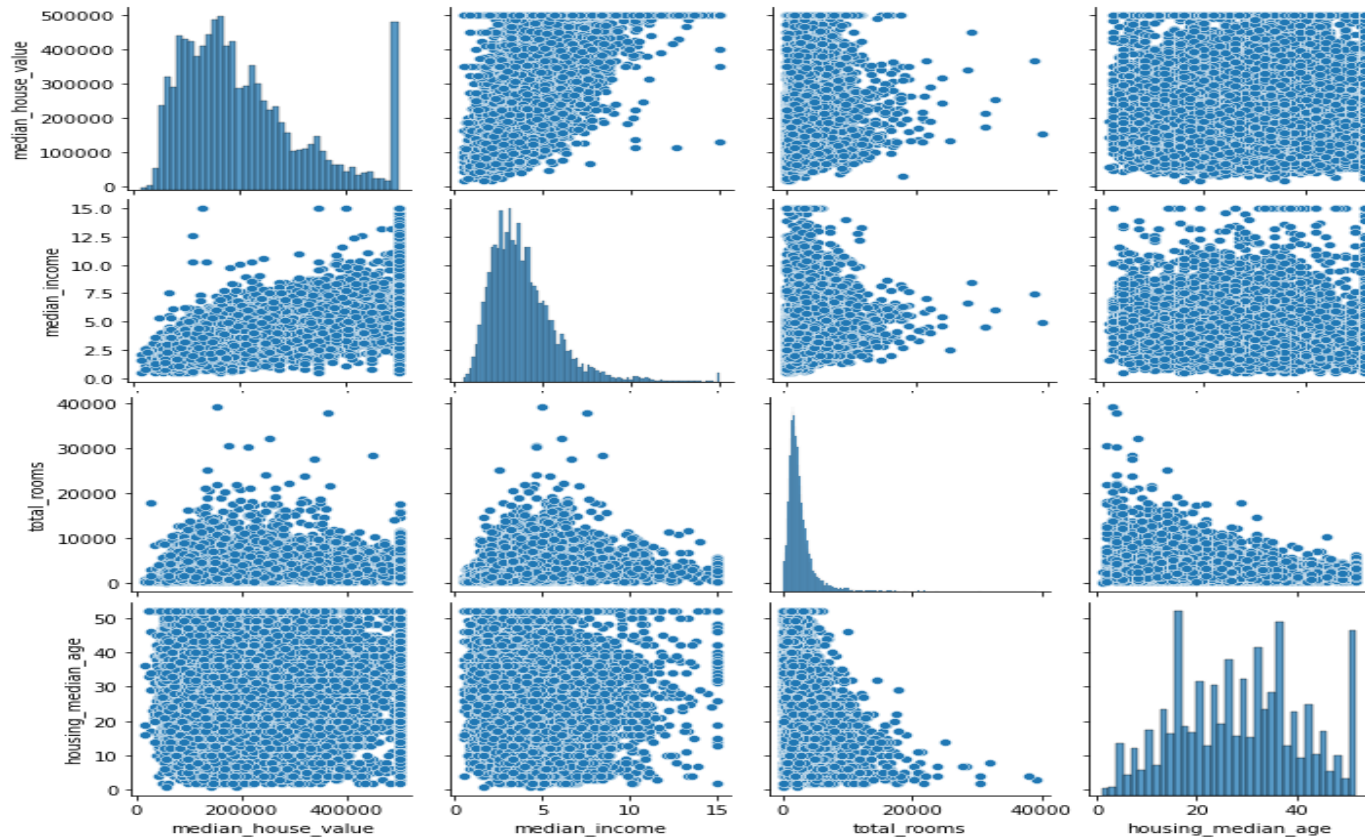
The above output shows the correlation of median_house_value with other columns.

(Contd.)

Python execution O/P Screenshots and Interpretation

```
In [24]: sns.pairplot(data[['median_house_value', 'median_income', 'total_rooms', 'housing_median_age']])
```

Out[24]: <seaborn.axisgrid.PairGrid at 0xca38640>



Median Income is the most promising attribute to get Median Housing Price.


(Contd.)

Python execution O/P Screenshots and Interpretation

```
In [25]: plt.figure(figsize=(10,6))  
plt.scatter(y=data['median_house_value'],x=data['median_income'],alpha=0.3)
```



```
Out[25]: <matplotlib.collections.PathCollection at 0xfb605e0>
```



A clear line can be seen at 500k at which the data is capped. Similar lines can be seen around 450k,350k. This kind of data may degrade the performance of model.

(Contd.)

Python execution O/P Screenshots and Interpretation



Creating new features:

- rooms per household
- bedrooms per room
- population per household


Initially creating them on a copy of dataset and checking whether they are making any difference.

```
In [26]: data1=data.copy()
```

```
In [27]: data1['rooms_per_household']=data1['total_rooms']/data1['households']  
data1['bedrooms_per_room']=data1['total_bedrooms']/data1['total_rooms']  
data1['population_per_household']=data1['population']/data1['households']
```

```
In [28]: data1.head(1)
```

Out[28]:



	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	ocean_proximity	median_house_value	room
0	-121.89	37.29	38	1568	351.0	710	339	2.7042	<1H OCEAN	286600	

The above output shows the first row of “data1” dataset.

(Contd.)

Python execution O/P Screenshots and Interpretation

```
In [29]: corr_matrix=data1.corr()
```

```
In [30]: corr_matrix['median_house_value'].sort_values(ascending=False)
```



```
Out[30]: median_house_value      1.000000  
         median_income        0.687160  
         rooms_per_household    0.146285  
         total_rooms           0.135097  
         housing_median_age     0.114110  
         households            0.064506  
         total_bedrooms         0.047689  
         population_per_household -0.021985  
         population            -0.026920  
         longitude             -0.047432  
         latitude              -0.142724  
         bedrooms_per_room      -0.259984  
         Name: median_house_value, dtype: float64
```

It is clear that 'rooms_per_household' and 'bedrooms_per_room' have better correlation with 'median_house_value' than 'total_rooms' and 'total_bedrooms'.

(Contd.)

Python execution O/P Screenshots and Interpretation



I have created classes for individual data manipulation job, so that I can add them into a pipeline.

```
In [31]: class FeaturesAdder(BaseEstimator,TransformerMixin):
        """This class adds new features in the dataset.
           Features added are : rooms_per_household, bedrooms_per_room, and population_per_household.
        """
        def fit(self,X,y=None):
            return self
        def transform(self,X,y=None):
            X['rooms_per_household']=X['total_rooms']/X['households']
            X['bedrooms_per_room']=X['total_bedrooms']/X['total_rooms']
            X['population_per_household']=X['population']/X['households']
            return X
```




Null Value Imputation

```
In [32]: #Get number of Null Values
        def get_null_count(data):
            for i in data.columns:
                print(i,': ',len(data[data[i].isnull()][i]))
```

(Contd.)

Python execution O/P Screenshots and Interpretation

```
In [33]: get_null_count(data)
```



```
longitude : 0  
latitude : 0  
housing_median_age : 0  
total_rooms : 0  
total_bedrooms : 158  
population : 0  
households : 0  
median_income : 0  
ocean_proximity : 0  
median_house_value : 0
```

Here, we can see in the output that, only “total_bedrooms” contains null values and the total number of null values are 158.

```
In [34]: data_null=data[data['total_bedrooms'].isnull()]
```



Plotting null values to check for any patterns.

For example: whether all nulls are concentrated in region or not.


(Contd.)

Python execution O/P Screenshots and Interpretation

```
In [35]: plt.figure(figsize=(12,6))  
plt.scatter(x=data_null['longitude'],y=data_null['latitude'],alpha=0.5, label='latitude')  
plt.legend()
```



Out[35]: <matplotlib.legend.Legend at 0xfb9a790>



We can see from the above plot that, there is not any specific pattern in null values except a dense spot near -118 longitude.

(Contd.)

Python execution O/P Screenshots and Interpretation

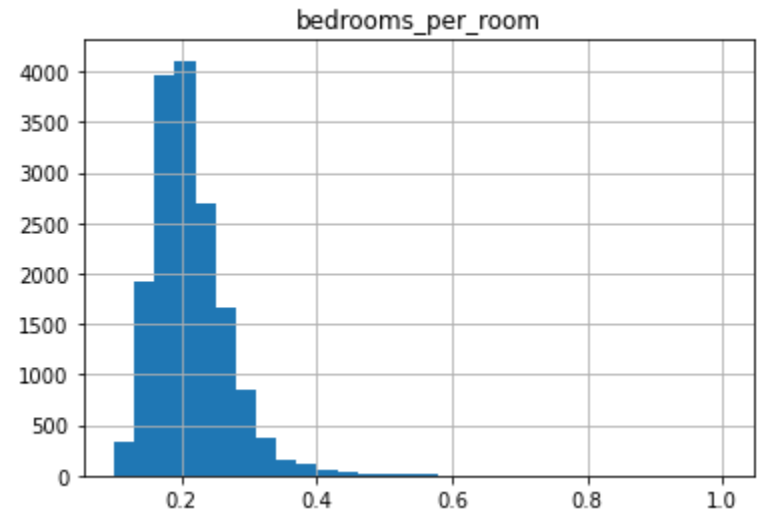



Let's plot histograms of columns with null values.

```
In [36]: data1.hist(column='total_bedrooms',bins=30)
data1.hist(column='bedrooms_per_room',bins=30)
```



```
Out[36]: array([[<AxesSubplot:title={'center':'bedrooms_per_room'}>]] dtype=object)
```



As the distributions of 'total_bedrooms' and 'bedrooms_per_room' are skewed, it's better to replace null values by median. I have used sklearn's Imputer to do this job.

(Contd.)


Python execution O/P Screenshots and Interpretation



Removing Outliers



```
In [37]: num_features=['longitude', 'latitude', 'housing_median_age', 'total_rooms',  
    'total_bedrooms', 'population', 'households', 'median_income',  
    'median_house_value', 'rooms_per_household',  
    'bedrooms_per_room', 'population_per_household']
```

```
In [38]: for i in num_features:  
    fig, ax = plt.subplots()  
    fig.set_size_inches(12,6)  
    sns.boxplot(x=i,data=data1,ax=ax)
```





(Contd.)

Python execution O/P Screenshots and Interpretation





(Contd.)

Python execution O/P Screenshots and Interpretation





(Contd.)

Python execution O/P Screenshots and Interpretation





(Contd.)

Python execution O/P Screenshots and Interpretation




(Contd.)

Python execution O/P Screenshots and Interpretation



(Contd.)

Python execution O/P Screenshots and Interpretation



By studying above box plots in details as well as scatter plot of 'Median_House_Value' vs. 'Median_Income', I decided following conditions to remove outliers.

(Contd.)

Python execution O/P Screenshots and Interpretation

```
In [39]: class RemoveOutliers(BaseEstimator,TransformerMixin):
        """This class removes outliers from data.
        Note: Outlier values are hard coded
        """
        def fit (self,X,y=None):
            return self

        def transform(self,X,y=None):
            X=X[(X['median_house_value']!=500001) | (X['median_income']>=2)].reset_index(drop=True)
            X=X[X['median_income']<=11].reset_index(drop=True)
            X=X[(X['median_house_value']!=350000) | (X['median_income']>=1.5)].reset_index(drop=True)
            X=X[(X['median_house_value']!=450000) | (X['median_income']>=2)].reset_index(drop=True)
            X=X[(X['median_house_value']>=350000) | (X['median_income']<=9.5)].reset_index(drop=True)
            X=X[X['population']<=9000]
            X=X[(X['population_per_household']>=1.15) & (X['population_per_household']<=6.5)]
            X=X[X['rooms_per_household']<20]
            X=X[X['bedrooms_per_room']<0.5].reset_index(drop=True)
            return X
```

```
In [40]: data1=RemoveOutliers().fit_transform(data1)
```

```
In [41]: data_labels=data1['median_house_value']
        data1=data1.drop('median_house_value',axis=1)
```

(Contd.)




Python execution O/P Screenshots and Interpretation

➡ Transforming skewed features

```
In [42]: data1.hist(bins=50,figsize=(16,12))
```

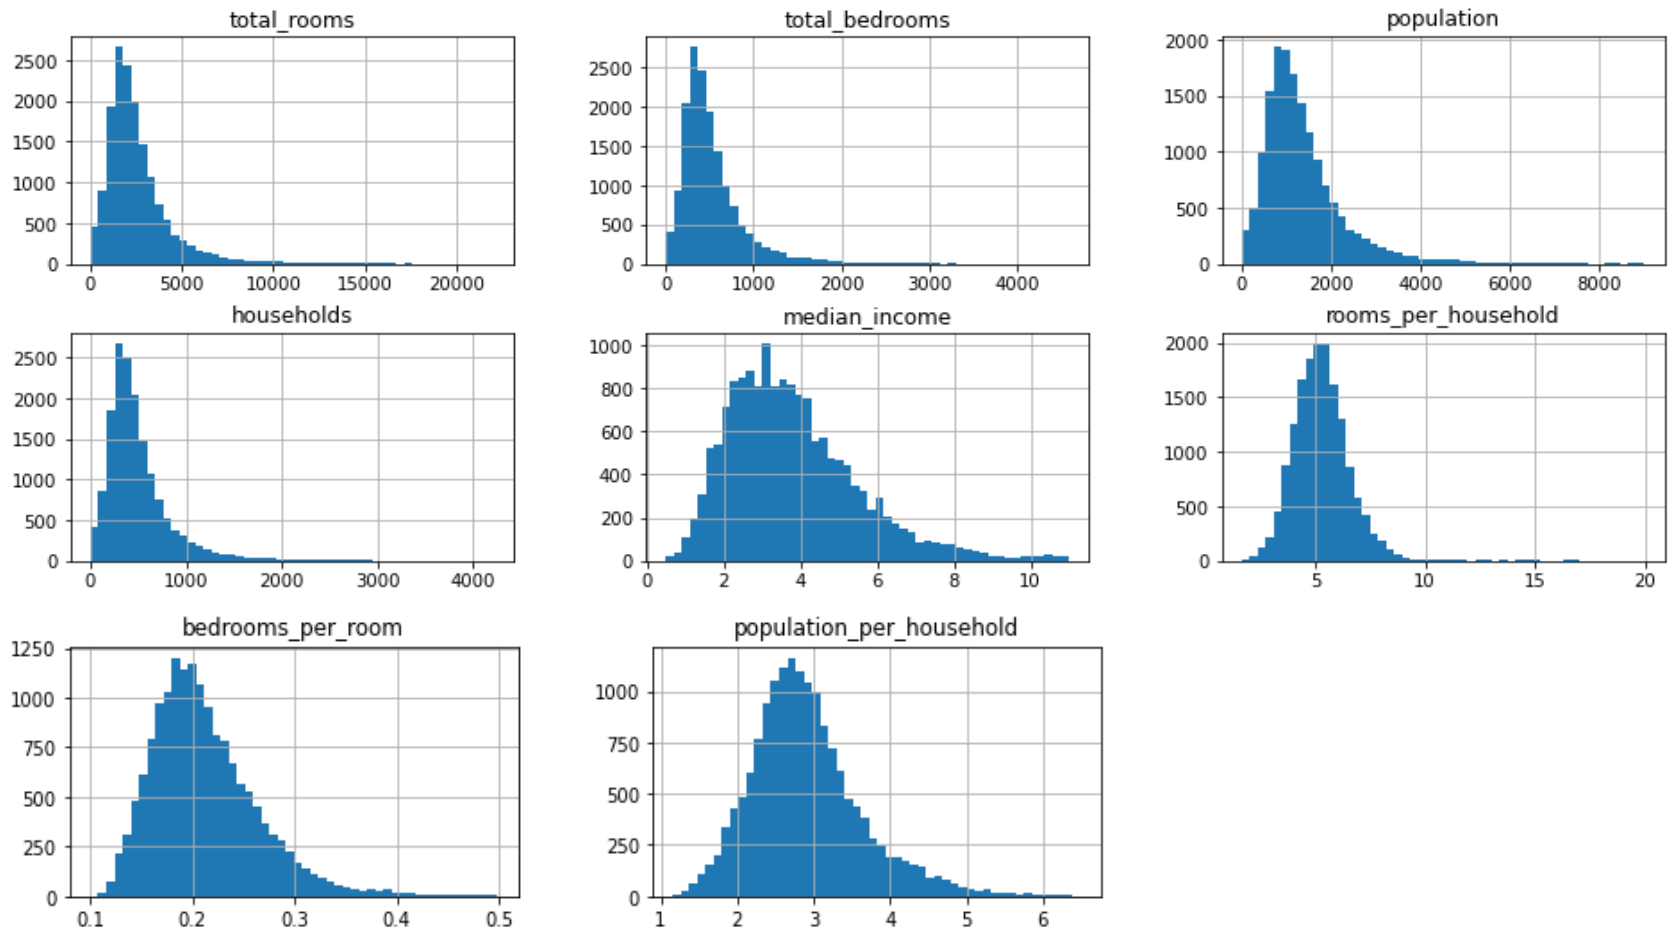
➡

```
Out[42]: array([[<AxesSubplot:title={'center':'longitude'}>,  
  <AxesSubplot:title={'center':'latitude'}>,  
  <AxesSubplot:title={'center':'housing_median_age'}>],  
  [<AxesSubplot:title={'center':'total_rooms'}>,  
  <AxesSubplot:title={'center':'total_bedrooms'}>,  
  <AxesSubplot:title={'center':'population'}>],  
  [<AxesSubplot:title={'center':'households'}>,  
  <AxesSubplot:title={'center':'median_income'}>,  
  <AxesSubplot:title={'center':'rooms_per_household'}>],  
  [<AxesSubplot:title={'center':'bedrooms_per_room'}>,  
  <AxesSubplot:title={'center':'population_per_household'}>],  
  <AxesSubplot:>]], dtype=object)
```



(Contd.)

Python execution O/P Screenshots and Interpretation



We can see from the plots that, after removing outliers, the skewness decreased.

(Contd.)

Python execution O/P Screenshots and Interpretation

```
In [43]: num_features=['longitude', 'latitude', 'housing_median_age', 'total_rooms',  
                    'total_bedrooms', 'population', 'households', 'median_income',  
                    'rooms_per_household',  
                    'bedrooms_per_room', 'population_per_household']
```



Get skewness of features

```
In [44]: skewness=[]  
for i in num_features:  
    skewness.append(data1[i].skew())  
pd.DataFrame(data=skewness,index=num_features,columns=['skewness']).sort_values(by='skewness',ascending=False)
```



Out[44]:

	skewness
total_rooms	2.902649
total_bedrooms	2.684070
households	2.629067
population	2.453927
rooms_per_household	2.168956
bedrooms_per_room	1.158469
median_income	1.051412
population_per_household	0.902046
latitude	0.462198
housing_median_age	0.061121
longitude	-0.293169

The above output shows the skewness of the attributes.

(Contd.)

Python execution O/P Screenshots and Interpretation




I transformed features using log1p and boxcox1p functions. boxcox1p is used for highly skewed features.

```
In [45]: data1['households'].apply(lambda x: np.log1p(x)**2).hist(bins=50,label='a')  
plt.legend()
```



Out[45]: <matplotlib.legend.Legend at 0x1254be50>



(Contd.)

Python execution O/P Screenshots and Interpretation

```
In [46]: class FeaturesTransformer(BaseEstimator,TransformerMixin):
        """This class trnsforms numerical featuress in the dataset.
        Note: Transformations are hard coded.
        """
        def fit(self,X,y=None):
            return self
        def transform(self,X,y=None):
            import numpy as np
            from scipy.special import boxcox1p
            X['total_rooms']=X['total_rooms'].apply(lambda x: boxcox1p(x,0.25))
            X['total_bedrooms']=X['total_bedrooms'].apply(lambda x: boxcox1p(x,0.25))
            X['households']=X['households'].apply(lambda x: boxcox1p(x,0.2))
            X['population']=X['population'].apply(lambda x: boxcox1p(x,0.3))
            X['rooms_per_household']=X['rooms_per_household'].apply(lambda x: np.log1p(x)**0.5)
            X['bedrooms_per_room']=X['bedrooms_per_room'].apply(lambda x: np.log1p(x)**0.25)
            X['median_income']=X['median_income'].apply(lambda x: np.log1p(x)**1.25)
            X['population_per_household']=X['population_per_household'].apply(lambda x: np.log1p(x)**1)
            return X
```

```
In [47]: data1=FeaturesTransformer().fit_transform(data1)
```

(Contd.)

Python execution O/P Screenshots and Interpretation






Transformed features

```
In [48]: data1.hist(bins=50,figsize=(15,12))
```

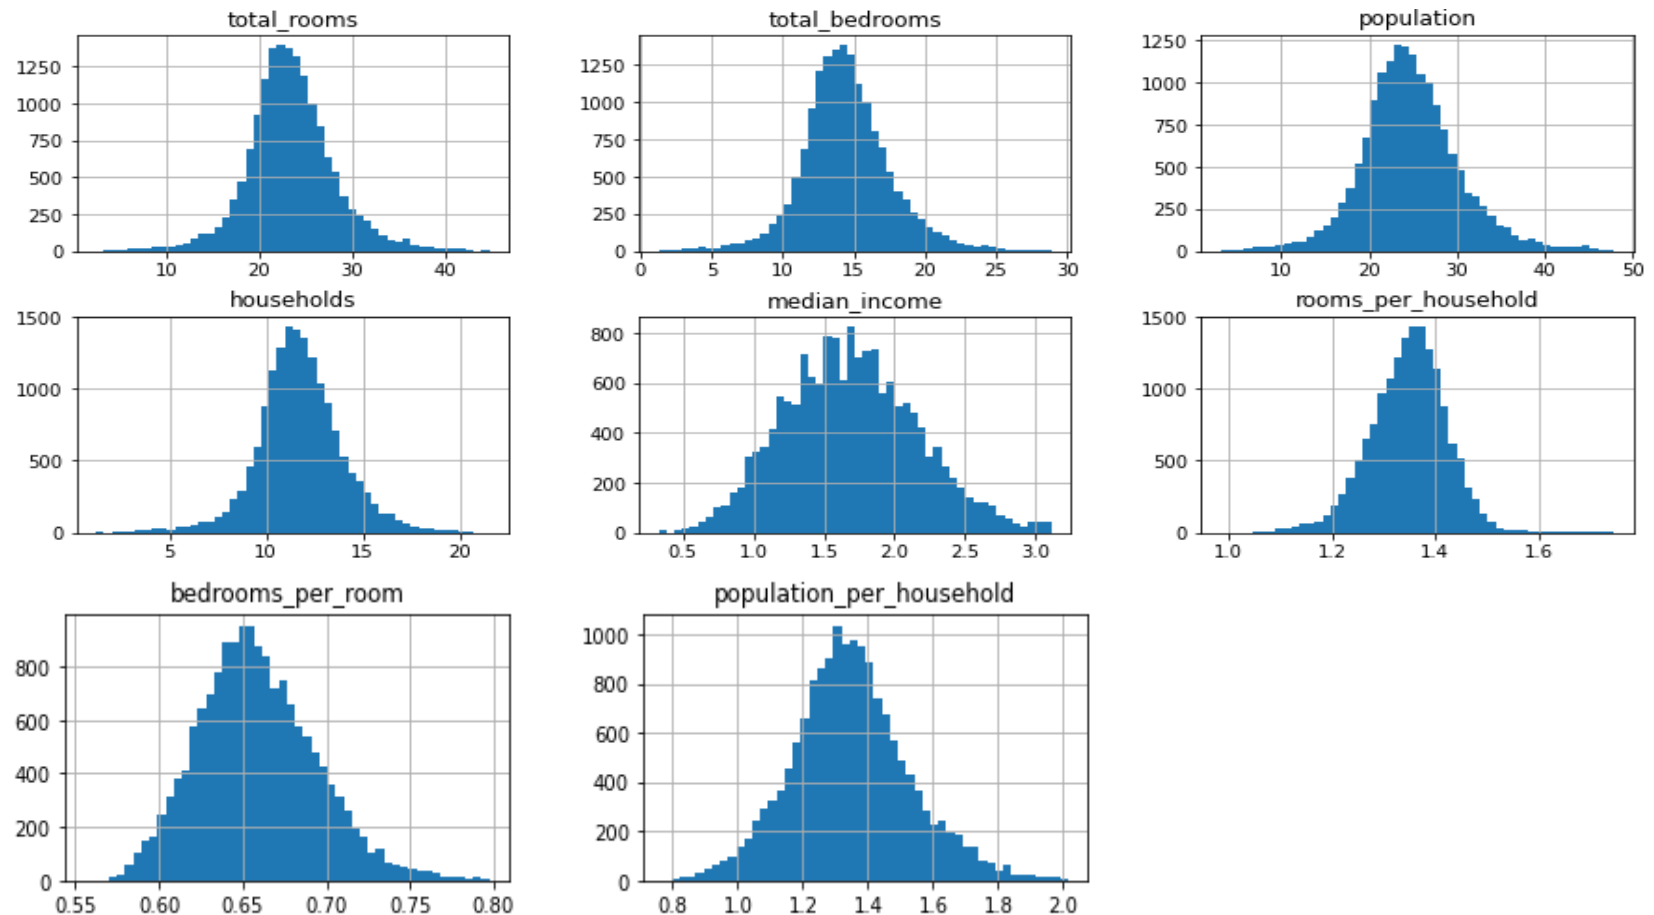


```
Out[48]: array([[<AxesSubplot:title={'center':'longitude'}>,  
                <AxesSubplot:title={'center':'latitude'}>,  
                <AxesSubplot:title={'center':'housing_median_age'}>],  
               [<AxesSubplot:title={'center':'total_rooms'}>,  
                <AxesSubplot:title={'center':'total_bedrooms'}>,  
                <AxesSubplot:title={'center':'population'}>],  
               [<AxesSubplot:title={'center':'households'}>,  
                <AxesSubplot:title={'center':'median_income'}>,  
                <AxesSubplot:title={'center':'rooms_per_household'}>],  
               [<AxesSubplot:title={'center':'bedrooms_per_room'}>,  
                <AxesSubplot:title={'center':'population_per_household'}>],  
               <AxesSubplot:>]], dtype=object)
```



(Contd.)

Python execution O/P Screenshots and Interpretation



We can see from the plots that, after feature transform, distribution become normal.

(Contd.)

Python execution O/P Screenshots and Interpretation



Getting dummy variables and Feature Scaling

```
In [49]: from sklearn.pipeline import Pipeline, FeatureUnion
         from sklearn.preprocessing import StandardScaler
         from sklearn.impute import SimpleImputer
```

```
In [50]: imputer = SimpleImputer(missing_values=np.nan, strategy='mean')
```

```
In [51]: class DataFrameSelector(BaseEstimator, TransformerMixin):
         """This class is a dataframe selector.
         Data members:
             features: A list of column_names you want in output dataframe
         """
         def __init__(self, features):
             self.features = features
         def fit(self, X, y=None):
             return self
         def transform(self, X, y=None):
             return X[self.features]
```

```
In [52]: class GetDummies(BaseEstimator, TransformerMixin):
         """This class is used to get dummy columns from categorical columns."""
         def fit(self, X, y=None):
             return self
         def transform(self, X, y=None):
             #change ISLAND to NEAR BAY...as count of ISLAND is very low
             X[X=='ISLAND']='NEAR BAY'
             return (pd.get_dummies(X, drop_first=True))
```

```
In [53]: num_features=['longitude', 'latitude', 'housing_median_age', 'total_rooms',
                       'total_bedrooms', 'population', 'households', 'median_income',
                       'rooms_per_household', 'bedrooms_per_room', 'population_per_household']
         cat_features=['ocean_proximity']
```

(Contd.)

Python execution O/P Screenshots and Interpretation

```
In [54]: num_pipeline=Pipeline([
          ('selector',DataFrameSelector(num_features)),
          ('imputer',SimpleImputer(strategy='mean')),
          ('std_scaler',StandardScaler())
        ])

#Preparing categorical data
cat_pipeline=Pipeline([
          ('selector',DataFrameSelector(cat_features)),
          ('get_dummies',GetDummies())
        ])

#Combining numerical and categorica data
data_cleaning_pipeline=FeatureUnion(transformer_list=[
          ('num_pipeline',num_pipeline),
          ('cat_pipeline',cat_pipeline),
        ])
])
```

```
In [55]: data1=data_cleaning_pipeline.fit_transform(data1)
```



Above pipeline returned an array. I converted it to data frame again so that we can look at columns easily and do further processing.

```
In [56]: l=num_features.copy()
          l.extend([0,1,2])
```


```
In [57]: data1=pd.DataFrame(data1,columns=l)
```

(Contd.)

Python execution O/P Screenshots and Interpretation

```
In [58]: data1.head(2)
```

Out[58]:



	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	rooms_per_household	bedrooms_per_room
0	-1.154224	0.772567	0.746370	-0.475425	-0.369225	-0.743301	-0.314063	-0.611150	-0.449849	0.334068
1	-1.174180	0.660361	-1.171019	-1.491029	-1.797235	-1.693043	-1.732155	1.464804	0.622682	-1.106072



Scaling labels:

➤ This is necessary for some machine learning algorithms

```
In [59]: y=data_labels.copy()
```

```
In [60]: label_scaler=StandardScaler()  
data_labels=label_scaler.fit_transform(y.values.reshape(-1,1))
```



Check for multicollinearity:

➤ Here I have removed features having VIF (Variance Inflation Factor) greater than 5 and p-Values greater than 0.05 I have written following functions to do the job.

(Contd.)

Python execution O/P Screenshots and Interpretation

```
In [61]: def get_vif(X):
import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor
l = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
return pd.Series(index=X.columns,data=l).sort_values(ascending=False)

def remove_by_pvalue(X,y,pvalue=0.05):
    """Remove features with p-value more than 'pvalue'

    This function uses statsmodels.api.OLS model. Please add intercept to data externally.
    Input:
        X: Array or dataframe excluding predicted variable
        y: Series or list of predicted variable
        pvalue: int or float

    Note:
        X is changed inplace
    """
    import statsmodels.api as sm
    for i in range(len(X.columns)):
        regressor_OLS=sm.OLS(endog=y,exog=X).fit()
        s=regressor_OLS.pvalues.sort_values(ascending=False)
        if s.iloc[0]>pvalue:
            X.drop(s.index[0],axis=1,inplace=True)
            print('Removed: ',s.index[0],'P-value: ',s.iloc[0])

def remove_by_vif(X,vif=5):
    """Remove columns from X whose VIF is greater than supplied 'vif'
    Parameters:
        X:array or dataframe containing data excluding target variable
        vif: int or float of limiting value of VIF
    Note:
        This function changes X inplace
    """
    import statsmodels.api as sm
    from statsmodels.stats.outliers_influence import variance_inflation_factor

    for i in range(len(X.columns)):
        l = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
        s=pd.Series(index=X.columns,data=l).sort_values(ascending=False)
        if s.iloc[0]>vif:
            X.drop(s.index[0],axis=1,inplace=True)
            print('Removed: ',s.index[0],', VIF: ',s.iloc[0])
        else:
            break
```

(Contd.)

Python execution O/P Screenshots and Interpretation

```
In [62]: # Get VIFs of all features  
get_vif(data1)
```




```
Out[62]: total_bedrooms      558.923558  
households      389.655010  
total_rooms     287.321351  
population     141.389762  
bedrooms_per_room  66.260643  
rooms_per_household  50.549888  
population_per_household  19.548814  
latitude       18.347840  
longitude      17.347168  
median_income   3.027807  
0              2.022556  
1              1.449785  
housing_median_age  1.374288  
2              1.223067  
dtype: float64
```

The above output showing the VIF (Variance Inflation Factor) of attributes of “data1”.

(Contd.)

Python execution O/P Screenshots and Interpretation

```
In [63]: remove_by_vif(data1)
```



```
Removed: total_bedrooms , VIF: 558.9235579608114  
Removed: households , VIF: 228.1741365426593  
Removed: total_rooms , VIF: 108.34813434774779  
Removed: latitude , VIF: 18.337136881413034  
Removed: bedrooms_per_room , VIF: 5.618818797870981
```

So, above 5 columns were removed because of VIF constraint.

```
In [64]: remove_by_pvalue(data1,data_labels)
```

No columns were removed by p-Value constraint.

 **This is all the data manipulation I have done. Following is the summarization:**

- Adding new features
- Removing outliers
- Transforming skewed features
- Null value imputation
- Dummy variables for ocean_proximity
- Check for multicollinearity
- Standard scaling

(Contd.)

Python execution O/P Screenshots and Interpretation


■ Second Portion : Training machine learning algorithms

I will try to find the best suitable machine learning model to predict house prices.

● Making Train Data ready

```
In [65]: df=pd.read_csv('strat_train_set.csv')
```

```
In [66]: df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 16512 entries, 0 to 16511
Data columns (total 10 columns):
#   Column              Non-Null Count  Dtype
---  -
0   longitude            16512 non-null  float64
1   latitude             16512 non-null  float64
2   housing_median_age   16512 non-null  int64
3   total_rooms          16512 non-null  int64
4   total_bedrooms       16354 non-null  float64
5   population           16512 non-null  int64
6   households            16512 non-null  int64
7   median_income        16512 non-null  float64
8   ocean_proximity      16512 non-null  object
9   median_house_value   16512 non-null  int64
dtypes: float64(4), int64(5), object(1)
memory usage: 1.3+ MB
```

The info() gives the information of the “strat_train_set” data frame like no. of non-null values, data type and memory usage.

(Contd.)

Python execution O/P Screenshots and Interpretation

```
In [67]: #Adding features
df=FeaturesAdder().fit_transform(df)

num_features=['longitude', 'latitude', 'housing_median_age', 'total_rooms',
              'total_bedrooms', 'population', 'households', 'median_income',
              'median_house_value', 'rooms_per_household',
              'bedrooms_per_room', 'population_per_household']

#Removing outliers
df=RemoveOutliers().fit_transform(df)

data_labels=df['median_house_value']
df=df.drop('median_house_value',axis=1)

#Transforming features
df=FeaturesTransformer().fit_transform(df)

num_features=['longitude', 'latitude', 'housing_median_age', 'total_rooms',
              'total_bedrooms', 'population', 'households', 'median_income',
              'rooms_per_household', 'bedrooms_per_room', 'population_per_household']
cat_features=['ocean_proximity']

#Mean value imputation, feature scaling, getting dummy variables
num_pipeline=Pipeline([
    ('selector',DataFrameSelector(num_features)),
    ('imputer',SimpleImputer(strategy='mean')),
    ('std_scaler',StandardScaler())
])

cat_pipeline=Pipeline([
    ('selector',DataFrameSelector(cat_features)),
    ('get_dummies',GetDummies())
])

#Combining numerical and categorical data
data_cleaning_pipeline=FeatureUnion(transformer_list=[
    ('num_pipeline',num_pipeline),
    ('cat_pipeline',cat_pipeline),
])

df=data_cleaning_pipeline.fit_transform(df)

l=num_features.copy()
l.extend([0,1,2])
df=pd.DataFrame(df,columns=l)
```

(Contd.)

Python execution O/P Screenshots and Interpretation

```
#Check for multicollinearity
remove_by_vif(df)
remove_by_pvalue(df,data_labels)

data_prepared=df
y_train=data_labels.copy()

label_scaler=StandardScaler()
data_labels=label_scaler.fit_transform(y_train.values.reshape(-1,1))
```



```
Removed: total_bedrooms , VIF: 558.9235579608114
Removed: households , VIF: 228.1741365426593
Removed: total_rooms , VIF: 108.34813434774779
Removed: latitude , VIF: 18.337136881413034
Removed: bedrooms_per_room , VIF: 5.618818797870981
```


So, above 5 columns were removed because of VIF constraint.

(Contd.)

Python execution O/P Screenshots and Interpretation

In [68]: data_prepared

Out[68]:



	longitude	housing_median_age	population	median_income	rooms_per_household	population_per_household	0	1	2
0	-1.154224	0.746370	-0.743301	-0.611150	-0.449849	-1.207634	0.0	0.0	0.0
1	-1.174180	-1.171019	-1.693043	1.464804	0.622682	-0.204224	0.0	0.0	0.0
2	1.185638	0.187132	-0.375164	-0.492940	-0.825039	-1.331664	0.0	0.0	1.0
3	-0.016721	-0.292216	0.285178	-1.299468	0.057913	1.602935	1.0	0.0	0.0
4	0.492161	-0.931346	2.390008	-0.368150	-0.558206	0.282264	0.0	0.0	0.0
...
16000	0.721657	1.385500	-1.008524	0.771706	0.651982	-0.173482	1.0	0.0	0.0
16001	1.006033	0.906153	-0.210085	-1.132569	-0.440662	1.539436	1.0	0.0	0.0
16002	1.584762	-1.570475	0.892959	-0.203348	0.844114	-0.152823	1.0	0.0	0.0
16003	0.781526	0.187132	0.169212	0.294852	0.266387	1.238031	0.0	0.0	0.0
16004	-1.433610	1.864848	0.067323	-0.004051	-0.259595	-1.405587	0.0	1.0	0.0

16005 rows x 9 columns


The “data_prepared” data frame is created from “df” having 16005 rows and 9 columns.

(Contd.)

Python execution O/P Screenshots and Interpretation

```
In [69]: data_prepared.describe()
```

Out[69]:



	longitude	housing_median_age	population	median_income	rooms_per_household	population_per_household	0	1
count	1.600500e+04	1.600500e+04	1.600500e+04	1.600500e+04	1.600500e+04	1.600500e+04	16005.000000	16005.000000
mean	2.320016e-15	1.304382e-16	5.637521e-16	1.228842e-16	5.576989e-16	-8.838776e-17	0.319838	0.112277
std	1.000031e+00	1.000031e+00	1.000031e+00	1.000031e+00	1.000031e+00	1.000031e+00	0.466428	0.315717
min	-2.381529e+00	-2.209605e+00	-3.858035e+00	-2.883611e+00	-4.797942e+00	-3.198401e+00	0.000000	0.000000
25%	-1.109322e+00	-8.514543e-01	-6.030186e-01	-7.133766e-01	-6.151884e-01	-6.358845e-01	0.000000	0.000000
50%	5.320735e-01	2.734921e-02	-5.728918e-02	-2.813425e-02	5.124493e-02	-4.460040e-02	0.000000	0.000000
75%	7.815258e-01	6.664790e-01	5.488315e-01	6.611860e-01	6.324853e-01	5.870352e-01	1.000000	0.000000
max	2.627473e+00	1.864848e+00	4.123765e+00	3.009781e+00	5.234390e+00	3.701446e+00	1.000000	1.000000

The describe() shows some basic statistical details like count, mean, median, standard deviation, quartile values, percentile, minimum and maximum values of each attributes of the dataset.

(Contd.)

Python execution O/P Screenshots and Interpretation

```
In [70]: data_prepared.isnull().sum()
```



```
Out[70]: longitude          0  
housing_median_age        0  
population                0  
median_income             0  
rooms_per_household       0  
population_per_household  0  
0                          0  
1                          0  
2                          0  
dtype: int64
```

Here, we can see in the output that, no attribute contains null values.

(Contd.)

Python execution O/P Screenshots and Interpretation

● Choosing ML Algorithm

```
In [71]: from sklearn.metrics import mean_squared_error, r2_score
         from sklearn.model_selection import learning_curve, cross_val_score, validation_curve, train_test_split
         from sklearn.model_selection import GridSearchCV
```

```
In [72]: def plot_validation_curve(scores, param_range, param_name, scoring='r2'):
         """This function plot validation curve.

         Parameters:
             scores: scores obtained from validation_curve() method
             param_range: list of range of parameters passed as 'param_range' in validation_curve() method
             scoring: str
         """
         n = len(param_range)
         if scoring == 'r2':
             train_score = [scores[0][i].mean() for i in range(0, n)]
             test_score = [scores[1][i].mean() for i in range(0, n)]
         elif scoring == 'neg_mean_squared_error':
             train_score = [np.sqrt(-scores[0][i].mean()) for i in range(0, n)]
             test_score = [np.sqrt(-scores[1][i].mean()) for i in range(0, n)]

         fig = plt.figure(figsize=(8, 6))
         plt.plot(param_range, train_score, label='Train')
         plt.plot(param_range, test_score, label='Test')
         plt.xticks = param_range
         plt.title("Validation curve of {}".format(param_name), size=12)
         plt.legend()
```

(Contd.)

Python execution O/P Screenshots and Interpretation



Linear Regression

```
In [73]: from sklearn.linear_model import LinearRegression,Ridge
```

```
In [74]: lr=LinearRegression()  
scores=cross_val_score(lr,data_prepared,data_labels,n_jobs=-1,cv=5,scoring='r2')  
print('R2: ',np.sqrt(scores).mean())
```



R2: 0.7991862993882449

To calculate RMSE I have used 'train_y' which are unscaled labels.

```
In [75]: lr=LinearRegression()  
scores=cross_val_score(lr,data_prepared,y_train,n_jobs=-1,cv=5,scoring='neg_mean_squared_error')  
print('RMSE: ',np.sqrt(-scores).mean())
```



RMSE: 67960.78780025143

So with linear regression,

- R-squared=0.8
- RMSE=67960.78
- RMSE of above model is 67960 which means that there's an average error of \$67960 in prediction of house price.

(Contd.)

Python execution O/P Screenshots and Interpretation



Ridge Regression

```
In [76]: ridge=Ridge(alpha=0.1,random_state=42)
scores=cross_val_score(ridge,data_prepared,data_labels,n_jobs=-1,cv=5,scoring='r2')
print('R2: ',np.sqrt(scores).mean())
```



R2: 0.7991863159569764

(Contd.)


Python execution O/P Screenshots and Interpretation

```
In [77]: ridge=Ridge(random_state=42)
param_name='alpha'
param_range = [1,10,50,100,200,300,500,700,1000]

scoring='neg_mean_squared_error'
curve=validation_curve(ridge,data_prepared,y_train,scoring=scoring,cv=5,param_name=param_name,
                        param_range=param_range,n_jobs=-1)

n=len(param_range)
'''if scoring=='r2':
    train_score=[curve[0][i].mean() for i in range (0,n)]
    test_score=[curve[i][i].mean() for i in range (0,n)]
elif scoring=='neg_mean_squared_error':
    train_score=[np.sqrt(-curve[0][i].mean()) for i in range (0,n)]
    test_score=[np.sqrt(-curve[i][i].mean()) for i in range (0,n)]
...

plot_validation_curve(curve,param_range,param_name,scoring)
```



➤ From this plot It can be seen that Train and Test scores are almost after alpha=700.

➤ Even at alpha=0, difference between RMSE is not much, so there is no objectionable over-fitting. As alpha is increased, bias is increasing and so the RMSE.

➤ As linear regression is not over-fitting the model, using Ridge regression is not necessary.


(Contd.)

Python execution O/P Screenshots and Interpretation

Support Vector Regression

```
In [78]: from sklearn.svm import SVR
```


```
In [79]: svr=SVR(degree=2)
param_grid={
    'C':[1,10,50,100]
}
scoring='r2'
grid=GridSearchCV(svr,param_grid,scoring=scoring,n_jobs=-1,cv=3,verbose=3)
grid.fit(data_prepared,data_labels)
```

 Fitting 3 folds for each of 4 candidates, totalling 12 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 10 out of 12 | elapsed: 4.4min remaining: 53.0s
[Parallel(n_jobs=-1)]: Done 12 out of 12 | elapsed: 5.5min finished
```

```
Out[79]: GridSearchCV(cv=3, estimator=SVR(degree=2), n_jobs=-1,
    param_grid={'C': [1, 10, 50, 100]}, scoring='r2', verbose=3)
```

```
In [80]: print(grid.best_params_)
```

 {'C': 10}

(Contd.)

Python execution O/P Screenshots and Interpretation

```
In [81]: svr=SVR(degree=2,C=10)
scores=cross_val_score(svr,data_prepared,data_labels,n_jobs=-1,cv=5,scoring='r2')
print('R2: ',np.sqrt(scores).mean())
```

R2: 0.8787045011060872

So, with support vector machines, $R^2=0.878$.
Which is better than Linear Regression model.



Decision Tree

```
In [82]: from sklearn.tree import DecisionTreeRegressor
```

```
In [83]: dtr=DecisionTreeRegressor(random_state=42)
scores=cross_val_score(dtr,data_prepared,y_train,n_jobs=-1,cv=5,scoring='r2')
print('R2: ',np.sqrt(scores).mean())
```

R2: 0.7571871948332872

- So, basic decision tree model is giving $R^2=0.7571$.
- Then I did some parameter tuning so get best decision tree model.


(Contd.)

Python execution O/P Screenshots and Interpretation



max_depth

```
In [84]: dtr=DecisionTreeRegressor(random_state=42)
param_name='max_depth'
param_range = range(2,21,2)
scoring='neg_mean_squared_error'
curve=validation_curve(dtr,data_prepared,y_train,scoring=scoring,cv=5,param_name=param_name,
                        param_range=param_range,n_jobs=-1)
plot_validation_curve(curve,param_range,param_name,scoring)
```



➤ From above diagram it is clear that model starts overfitting heavily after max_depth=8.

➤ max_depth: 3 to 8


(Contd.)

Python execution O/P Screenshots and Interpretation



max_features

```
In [85]: dtr=DecisionTreeRegressor(max_depth=4,random_state=42)
param_name='max_features'
param_range = range(2,10)
scoring='neg_mean_squared_error'
curve=validation_curve(dtr,data_prepared,y_train,scoring=scoring,cv=5,param_name=param_name,
    param_range=param_range,n_jobs=-1)
plot_validation_curve(curve,param_range,param_name,scoring)
```



So, we get the value of max_features: 6 to 9.

(Contd.)

Python execution O/P Screenshots and Interpretation




min_samples_split

```
In [86]: dtr=DecisionTreeRegressor(random_state=42)
param_name='min_samples_split'
param_range = range(2,2003,100)
scoring='neg_mean_squared_error'
curve=validation_curve(dtr,data_prepared,y_train,scoring=scoring,cv=5,param_name=param_name,
    param_range=param_range,n_jobs=-1)
plot_validation_curve(curve,param_range,param_name,scoring)
```



Validation curve of min_samples_split



With increasing min_samples_split after 200, variance is reducing.


(Contd.)

Python execution O/P Screenshots and Interpretation



min_samples_leaf

```
In [87]: dtr=DecisionTreeRegressor(max_depth=4,random_state=42)
param_name='min_samples_leaf'
param_range = range(1,701,100)
scoring='neg_mean_squared_error'
curve=validation_curve(dtr,data_prepared,y_train,scoring=scoring,cv=5,param_name=param_name,
    param_range=param_range,n_jobs=-1)
plot_validation_curve(curve,param_range,param_name,scoring)
```



Increasing min_samples_leaf is neither giving good results nor helping to reduce overfitting.

(Contd.)

Python execution O/P Screenshots and Interpretation



Gradient Boosting Regression

I used this ensembling method to get the best model from decision trees.

```
In [88]: from sklearn.ensemble import GradientBoostingRegressor
```

```
In [89]: gbr=GradientBoostingRegressor(random_state=42)
param_grid={
    'n_estimators':[100,500,1000],
    'max_depth':range(3,9,1),
    'max_features':[6,7,9],
    'min_samples_split':[200,400]
}
scoring='r2'
grid=GridSearchCV(gbr,param_grid,scoring=scoring,n_jobs=-1,cv=3,verbose=1)
grid.fit(data_prepared,y_train)
print(grid.best_params_)
```



Fitting 3 folds for each of 108 candidates, totalling 324 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 42 tasks      | elapsed:  2.2min
[Parallel(n_jobs=-1)]: Done 192 tasks   | elapsed: 14.8min
[Parallel(n_jobs=-1)]: Done 324 out of 324 | elapsed: 34.0min finished
```


```
{'max_depth': 8, 'max_features': 6, 'min_samples_split': 200, 'n_estimators': 100}
```

(Contd.)

Python execution O/P Screenshots and Interpretation

➡ Check for n_estimators


```
In [90]: gbr=GradientBoostingRegressor(max_depth= 8, max_features=6, min_samples_split=200,random_state=42)
          param_name='n_estimators'
          param_range = range(50,251,50)
          scoring='r2'
          curve=validation_curve(gbr,data_prepared,y_train,scoring=scoring,cv=5,param_name=param_name,
                                param_range=param_range,n_jobs=-1)
          plot_validation_curve(curve,param_range,param_name,scoring)
```




- With increasing number of estimators, model is overfitting.
- With lower number of estimators, model is underfitting.

(Contd.)

Python execution O/P Screenshots and Interpretation

 I chose `n_estimators=100`

```
In [91]: gbr=GradientBoostingRegressor(max_depth= 8, max_features=6, min_samples_split=200, n_estimators=100,random_state=42)
scores=cross_val_score(gbr,data_prepared,y_train,n_jobs=-1,cv=5,scoring='r2')
print('R2: ',np.sqrt(scores).mean())
```

 R2: 0.8974790145435815

- With above Gradient Boosting model, R2=0.897.
- This model is even better than support vector regressor.

 **Stacking**

- This is another ensembling method.
- In this, I have used Linear regression and Gradient boosting regression as base models and Support Vector Regression as meta model.

(Contd.)

Python execution O/P Screenshots and Interpretation

```
In [92]: from mlxtend.regressor import StackingRegressor

In [93]: lr=LinearRegression()
svr=SVR(degree=2,C=10)
gbr=GradientBoostingRegressor(max_depth= 8, max_features=6, min_samples_split=200, n_estimators=100,random_state=42)

In [94]: sr=StackingRegressor([lr,gbr],svr,verbose=3)
scores=cross_val_score(sr,data_prepared,data_labels,n_jobs=-1,cv=3,scoring='r2')
print('R2: ',np.sqrt(scores).mean())
```

R2: 0.8920759836217309

For stacking,

- R2=0.892
- This is less than Gradient Boosting Regressor.

Final Model:

I finalized following Gradient Boosing Regressor with following parameters:

- n_estimators=100
- max_depth=8
- max_features=6
- min_samples_split=200

(Contd.)


Python execution O/P Screenshots and Interpretation

Learning Curve

```
In [95]: gbr=GradientBoostingRegressor(max_depth= 8, max_features=6, min_samples_split=200, n_estimators=100,random_state=42)
lc=learning_curve(gbr,data_prepared,y_train,cv=5,n_jobs=-1)
size=lc[0]
train_score=[lc[1][i].mean() for i in range (0,5)]
test_score=[lc[2][i].mean() for i in range (0,5)]
fig=plt.figure(figsize=(12,8))
plt.plot(size,train_score,label='Train')
plt.plot(size,test_score,label="Test")
plt.legend()
```



Out[95]: <matplotlib.legend.Legend at 0x15071400>



➤ From learning curve it is clear that this model is still overfitting. But, with the increase in number of samples, variance is decreasing.

➤ So, getting more training data can be a solution for overfitting.

(Contd.)

Python execution O/P Screenshots and Interpretation

Predicting on Test Data Set

```
In [96]: test_data=pd.read_csv('strat_test_set.csv')
test_data_labels=test_data['median_house_value'].copy()
test_data.drop('median_house_value',axis=1,inplace=True)
```



I did following data manipulations on test data

- Adding features
- Transforming skewed features

```
In [97]: fa=FeaturesAdder()
df_test=fa.fit_transform(test_data)

df_test=FeaturesTransformer().fit_transform(df_test)

#used 'transform' method, as I don't want to refit the standard scalar on test data.
df_test=data_cleaning_pipeline.transform(df_test)

df_test=pd.DataFrame(df_test,columns=1)
```

(Contd.)

Python execution O/P Screenshots and Interpretation

Scaling test data labels

In [98]: *#used 'transform' method, as I don't want to refit the standard scalar on test data.*

```
y_test=test_data_labels.copy()
test_data_labels=label_scaler.transform(y_test.values.reshape(-1,1))
```

Choosing same columns of test_data as of train_data

In [99]: `test_data_prepared=df_test[data_prepared.columns]`

In [100]: `gbr=GradientBoostingRegressor(max_depth= 8, max_features=6, min_samples_split=200, n_estimators=100,random_state=42)`
`gbr.fit(data_prepared,data_labels)`
`train_pred=gbr.predict(data_prepared)`
`test_pred=gbr.predict(test_data_prepared)`

In [101]: `mse=mean_squared_error(y_train,label_scaler.inverse_transform(train_pred))`
`rmse=np.sqrt(mse)`
`print('train error:',rmse)`
`mse=mean_squared_error(y_test,label_scaler.inverse_transform(test_pred))`
`rmse=np.sqrt(mse)`
`print('test error:',rmse)`



train error: 39773.94783299169
test error: 49795.96894045239

So, the train error : 39773.94783299169 and test error : 49795.96894045239.